

CS 140 Project 1: RSDL Variant Implementation

This documentation contains all the information regarding the project including but not limited to the list of changes made in the files/codes per phase as well as the implementation for the different requirements such as the schedlog, swapping of sets, and replenishment of quanta. This document also contains the link for the google drive which contains the document video to be passed.

Introduction

This project is similar to the lab exercise we did in the subject, mainly lab exercise 5 from which we got the schedlog code with minor changes as well as the idea for the project. Project 1 is mainly this exercise with added functions divided into 5 phases. Some of these functions are making the originally round-robin scheduler FIFO based, creation of two sets called active and expired which have levels, implementation for enqueueing and dequeuing of the processes as well as some rules for swapping of sets as well as replenishment of the quanta of both processes and level quantum. All of these will be discussed thoroughly throughout the project documentation.

Important Auxiliary Functions and Other Changes

There are several helper functions used throughout the project to help their implementation. Listed below are these functions as well as their code snippets and explanation.

rsdl.h

The file rsdl.h contains #define statements for constants that shall be used throughout the implementation of this project. Shown below is a code snippet that contains the said information:

```
1  #define RSDL_LEVELS 3
2  #define RSDL_STARTING_LEVEL 0
3  #define RSDL_PROC_QUANTUM 50
4  #define RSDL_LEVEL_QUANTUM 100
```

Figure i. Code Snippet of rsdl.h

The file defines the values of four constants. First is the `RSDL_LEVELS` which is set to 3. This corresponds to the number of priority levels as there are three levels for the active set and three levels for the expired set. The second constant defined is `RSDL_STARTING_LEVEL` which we have set to 0 since all processes will start at active level 0 (with the exception of those that called the `priofork` system call). The next constant defined is `RSDL_PROC_QUANTUM` which is the length of the quantum (in ticks) assigned by default to each process. We have set its value to 50 ticks. Lastly, the `RSDL_LEVEL_QUANTUM` is the length of the quantum (also in ticks) assigned by default to each level. We have set its value to 100 ticks.

queueclean

This helper function mainly cleans the queue for several factors. The code for this function is shown below. To do this, we first find the highest unused process. There are two parts and conditions for using `queueclean`. The first one is clearing up the queue in the active set after the processes goes on to the expired set. This allows the program to have the correct order in queue once the processes go back to the active set. To do this, we first check if the process is in the expired set as well as if this process is running and has no quantum left. If so, we use the for loop to clear the queue for all the processes that went to the expired set. Another time for using the `queueclean` is for when there are zombie and unused processes, processes that went down a level, as well as for finished processes. To do this, we once again check the conditionals and if so, we loop through to clean the queue.

```
412 void queueclean(void){
413     //Remove UNUSED, ZOMBIE, and PROC with 0 Quantum from the queue of both the active and expired set.
414     struct proc *pp;
415     struct proc *xp;
416     int temp_active = 0;
417     int highest_idx = -1;
418
419     //Identify highest idx
420     for (int k = 0; k < NPROC; k++) {
421         pp = &ptable.proc[k];
422         if (pp->state != UNUSED) highest_idx = k;
423     }
424     //Clear up Queue in the Active Set for processes done after a process went to expired
425     for (int k = 0; k <= highest_idx; k++) {
426         pp = &ptable.proc[k];
427         if(pp->state == RUNNING && pp->ticks_left == 0 && pp->set == 1 && pp->level == pp->starting_level){
428             temp_active = pp->queue;
429             pp->queue = -1;
430             for (int i = 0; i<=highest_idx; i++){
431                 xp = &ptable.proc[i];
432                 if(temp_active < xp->queue && xp->set == 0 && xp->level == RSDL_LEVELS-1)
433                     xp->queue = xp->queue - 1;
434             }
435         }
436     }
```

Figure ii. Snippet Code of `queueclean` Function

```
438 //Clear up Queue in the Active Set for process done but went down a level and for zombie and unused processes
439 for (int j = 0; j < RSDL_LEVELS; j++){
440     for (int k = 0; k <= highest_idx; k++) {
441         pp = &ptable.proc[k];
442         if((((pp->state == UNUSED && pp->queue != -1) || (pp->state == ZOMBIE && pp->queue != -1)) && pp->set == 0 && pp->level == j) ||
443             (pp->state == RUNNING && pp->ticks_left == 0 && pp->set == 0 && pp->level == j+1)){
444             temp_active = pp->queue;
445             pp->queue = -1;
446             for (int i = 0; i<=highest_idx; i++){
447                 xp = &ptable.proc[i];
448                 if(temp_active < xp->queue && xp->set == 0 && xp->level == j)
449                     xp->queue = xp->queue - 1;
450             }
451         }
452     }
453 }
454 }
```

Figure iii: Continuation of the Snippet Code of queueclean Function

queuesort

This is another helper function where its main purpose is to find the next queue which is concerned with the order of the processes. The code for this function is shown below. To do this, we first find the highest unused process as well as clean the queue using the auxiliary function queueclean. We then find the next queue by looping taking into account the sets and the levels. After finding the highest queue, we increment 1 to the value of it to get the value of the next queue.

```
482 int queuesort(int pset, int plevel){
483     //acquire(&ptable.lock);
484     //Find the next queue
485     struct proc *pp;
486     int nextqueue = -1;
487     int highest_idx = -1;
488
489     //Identify highest idx
490     for (int k = 0; k < NPROC; k++) {
491         pp = &ptable.proc[k];
492         if (pp->state != UNUSED) highest_idx = k;
493     }
494
495     //release(&ptable.lock);
496     //Call queueclean which cleans the queue of both active and expired sets
497     queueclean();
498     //UNUSED and ZOMBIE processes are clear. Find the highest in the queue
499     //Consider Set and Level
500     for (int k = 0; k <= highest_idx; k++) {
501         pp = &ptable.proc[k];
502         if((pp->state == SLEEPING || pp->state == RUNNING || pp->state == RUNNABLE) && pp->set == pset && pp->level == plevel){
503             if(pp->queue > nextqueue)
504                 nextqueue = pp->queue;
505         }
506     }
507     return nextqueue+1;
508 }
```

Figure iv: Code Snippet of queuesort Function

expired_check

This helper function's main purpose is to check if there is a process left in the specific level (set by the user) in the expired set. This works by looping through the maximum amount of processes (64) and returning 1 if there is a process found. We know this if the value of the set of the process is 1 and its level is the same as the parameter.

```
int expired_check(int lvl){
    struct proc *xp;
    for(int k = 0; k < NPROC; k++){
        xp = &ptable.proc[k];
        //Find if there is a process in the expired level set
        if(xp->set == 1 && xp->level == lvl) return 1;
    }
    return 0;
}
```

Figure v: Code for Expired_Check Function

active_check

Similar to the auxiliary function expired_check, this function checks on whether there are functions in the active set in the specific level set by the user. The conditionals for this are a bit different however since we want processes that are initialized and not killed. We also added the condition for the state of the process found to not be unused since this means they weren't initialized yet and that they are not zombies.

```
int active_check(int lvl){
    struct proc *xp;
    for(int k = 0; k < NPROC; k++){
        xp = &ptable.proc[k];
        //Find if there is a process in the active level set
        if(xp->set == 0 && xp->level == lvl && xp->pid != 0 && xp->state != UNUSED && xp->state != ZOMBIE) return 1;
    }
    return 0;
}
```

Figure vi: Code Snippet for Auxiliary Function active_check

sleepclean

Similar to the auxiliary function queueclean, this function searches for the passed process coming from the sleep function (identified by its pid, and sleeping state) and

effectively rearrange the queue order as the process who went to sleep will be reenqueued to the back of the same level. This would place that process at the back.

```
456 void sleepclean(int spid){
457     struct proc *pp;
458     struct proc *xp;
459     int temp = 0;
460     int level = 0;
461     int highest_idx = -1;
462     for (int k = 0; k < NPROC; k++) {
463         pp = &ptable.proc[k];
464         if (pp->state != UNUSED) highest_idx = k;
465     }
466
467     for (int k = 0; k <= highest_idx; k++) {
468         pp = &ptable.proc[k];
469         if(pp->state == RUNNING && pp->pid == spid){
470             temp = pp->queue;
471             level = pp->level;
472             pp->queue = -1;
473             for (int i = 0; i<=highest_idx; i++){
474                 xp = &ptable.proc[i];
475                 if(temp < xp->queue && xp->set == 0 && xp->level == level)
476                     xp->queue = xp->queue - 1;
477             }
478         }
479     }
480 }
```

Figure vii. Code Snippet for Auxiliary Function sleepclean

Note that along with this auxiliary function, we modified the sleep() call by adding additional lines to accommodate this edge case. In lines 860 to 865, we added the sleepclean function and the queuesort to accommodate the reenqueue of the process involve in the sleep process.

```
858 // Go to sleep.
859 p->chan = chan;
860 sleepclean(p->pid);
861 p->state = SLEEPING;
862
863 release(&ptable.lock);
864 p->queue = queuesort(p->set, p->level);
865 acquire(&ptable.lock);
866
867 sched();
868
869 // Tidy up.
870 p->chan = 0;
```

Figure viii. Changes Made in the Sleep Syscall

1. List of all Phases with Working Code

In total, there are 5 phases of the project all of which have added functions over the previous one. For this project, the group was able to make up to the last phase of the project, **phase 5**. Codes for the different phases were saved in different branches. Phase 1 was saved in the branch **phase1**, phase 2 was saved in the branch **phase2**, phase 3 was saved in the branch **phase3**, phase 4 was saved in the branch **phase4**, and lastly phase 5 was saved in the branch **phase5**. For this project, we will discuss the entirety of the project meaning phase 5 only. The discussion of the last phase should include the discussions of the lower phases as phase 5 encompasses all the phases.

2. All references used with purpose specified

There were **no external references** used for any of the phases done in this project. The only reference used was the lab exercise 5 we were given in the past which used a schedlog code as well as the user test code (test.c). The group also made it as the basis for the project.

3. List of All Global Variables Introduced

There is only one global variable added in during the several phases of the project. These variables are the following:

- **Struct plevel** - used for storing the necessary values for each level such as the level, set, and the level quantum as well as for implementing the levels in phases 2 - 5. Struct lvltable was used for this and will be further discussed in the implementation of the levels.

4. List of Changes Made to the Process Control Block

There are a total of five (5) additions to the process control block which are:

- **ticks_left** - which was used for the amount of quantum/ticks left needed before a process finishes.
- **queue** - used for the ordering or queuing of processes in the right order
- **set** - used for recording/manipulating the active and expired sets.
- **level** - used for recording/finding the values in the different levels of the program (i.e. level 0 - N-1)
- **starting_level** - used for storing the value of the starting level of the processes.

5. Code Representation of the Active Set and its Levels

The code representation of the active set as well as its corresponding levels can be found in a function called **levelinit** in the `proc.c` file. The code involved in this is seen in figure 1 below.

```
void levelinit(void){
    struct qlevel *rp;
    //Init active levels
    for(int i=0;i<RSDL_LEVELS;i++){
        rp = &plevel.qlevel[i];
        rp->set = 0;
        rp->level = i;
        rp->lvlquantum = RSDL_LEVEL_QUANTUM;
    }
}
```

Figure 1: Code Snippet for Active Set and its Levels

Before we discuss this, we must first discuss its corequisite found in the file `proc.h` which has a new addition of a struct called `qlevel`. This contains the necessary value for initializing the set and its levels. The code for this is shown below.

```
65  struct qlevel{
66      int level;
67      int set;
68      int lvlquantum;
69  };
```

Figure 2: `qlevel` Struct Added in `Proc.h`

To discuss the code in figure 1, we first initialize the variable `*rp` with the struct `qlevel` which means each `rp` will hold the value for `level`, `set`, and `lvlquantum`. The group then used a for loop which loops through number 0 to N-1 for the levels initialization. This allowed the program to have their own values stored in a struct for each level in the active set.

6. Implementation Explanation for Initial Enqueuing of New Processes

The implementation for the initial enqueueing of a new process is very similar to the initial enqueueing of a new process done in the lab exercise 5 wherein if we found an unused process in the process table then we will initialize its needed values as well as change its state to embryo. The codes for this can be found in the **allocproc** function in `proc.c` file. Since there were additional functions in the phases of the project then there will be additions to the initial enqueueing implementation as well. These changes are shown and further discussed below.

To start with, the general explanation for the implementation of the initial enqueueing is to look for an unused process, if found change its state to embryo (a new process) and initialize its values, run in the kernel else return 0. There were no changes to the allocation of a stack to the process since the changes only mainly involved the initial values of the processes since there were levels and quantum added. These changes can be found in the **found** section of the `allocproc` function.

```
90  found:
91      p->state = EMBRYO;
92      p->pid = nextpid++;
93      p->set = 0;
94
95      //Check if starting level is available
96      int result = searchlevel(RSDL_STARTING_LEVEL);
97      if(result != -1)
98          p->level = result;
99      else{
100          p->set = 1;
101          p->level = RSDL_STARTING_LEVEL;
102      }
103
104      p->starting_level = RSDL_STARTING_LEVEL;
105      p->ticks_left = RSDL_PROC_QUANTUM;
106      release(&ptable.lock);
107
108      temp = queuesort(p->set, p->level);
109
110      acquire(&ptable.lock);
111      p->queue = temp;
112      release(&ptable.lock);
```

Figure 3: Changes Done in `allocproc` Function for New Initial Enqueueing Implementation

Figure 3 shows the relevant additions to the initialization of a new process. If there is an unused process found, we proceed as normal in making its state into an embryo (new process) as well as add a unique pid to it. The change however is that we also add a value to its `p->set` to 0 since we want the new processes to go to the active set instead of the expired set if they are new. We also search for the lowest active level using the function **searchlevel**. Figure 4 shows the code for the function search level and what it mainly does is search for the available active levels meaning active levels whose quantum are not yet depleted. This was done since there can be a possibility that the starting level has no available quantum already when the new process gets enqueued so the process doesn't start in the starting level or that no such active level exists so it starts out in the starting level in the expired set. After setting the new process' current level to be in any active level or in the starting level of the expired set, we also initialize other important values such as its `starting_level` which will be used for the swapping as well as its current quantum using the variable `ticks_left`. We then get its queue using the `queuesort` function which both finds the next queue as well as rid of the queue of zombie processes, finished processes, and downleveled processes and other more. We do this since we follow a FIFO based round-robin scheduler which means that all new processes must be at the last of the queue.

```
400 //Search available active levels
401 int searchlevel(int start){
402     struct qllevel *rp;
403     for(int i=0;i<RSDL_LEVELS;i++){
404         rp = &plevel.qllevel[i];
405         if(rp->level == start && rp->lvlquantum != 0) return start;
406         else if(rp->level > start && rp->lvlquantum != 0) return rp->level;
407     }
408     //Return -1 if there are no available active levels
409     return -1;
410 }
```

Figure 4: Code Snippet for searchlevel Function

7. Implementation Explanation for Dequeuing of Exiting Processes

The implementation for the dequeuing of exiting/finished processes is very simple. To account for the finished process we mainly use a part of the code in the **queueclean** function shown below. Notice how this code is used for both exiting processes as well as zombie and unused processes. The particular condition used for exiting processes is the line "**if pp->state == UNUSED && pp->queue != -1**" which means that a process only became unused since it exited successfully and we know this since its queue did not change. Notice in line 441 that we have the value of the queue of exiting processes to be

-1. This is how we know that the process we checked for has finished successfully and to dequeue this we need to change its queue value to -1. Doing so makes the program ignore this process in printing and other functions since it's no longer part of the queue.

```
438 //Clear up Queue in the Active Set for process done but went down a level and for zombie and unused processes
439 for (int j = 0; j < RSDL_LEVELS; j++){
440     for (int k = 0; k <= highest_idx; k++) {
441         pp = &ptable.proc[k];
442         if(((pp->state == UNUSED && pp->queue != -1) || (pp->state == ZOMBIE && pp->queue != -1)) && pp->set == 0 && pp->level == j) ||
443             (pp->state == RUNNING && pp->ticks_left == 0 && pp->set == 0 && pp->level == j+1)){
444             temp_active = pp->queue;
445             pp->queue = -1;
446             for (int i = 0; i<=highest_idx; i++){
447                 xp = &ptable.proc[i];
448                 if(temp_active < xp->queue && xp->set == 0 && xp->level == j)
449                     xp->queue = xp->queue - 1;
450             }
451         }
452     }
453 }
454 }
```

Figure 5: Code Snippet for Implementation of Dequeuing of Exiting Processes

8. Implementation Explanation for Process-Local Quantum Consumption

Relevant codes for Process-local quantum consumption can be found in the trap.c file along with Level-local quantum consumption. Code snippet regarding this is shown in figure 6 below. What the code in the snippet does is it checks whether there is a process and whether it is running as well as check if its trapframe trap number is the same as the previous user space registers before its context switches. The important part here however is the code `--myproc()->ticks_left` which is the main code for the quantum consumption process-wise. What it does is for every tick, it decrements its quantum (the ticks_left variable) by 1.

```
if(myproc() && myproc()->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER){
    for(rp = plevel.qlevel; rp < &plevel.qlevel[RSDL_LEVELS*2]; rp++){
        if(rp->level == myproc()->level && rp->set == myproc()->set){
            //cprintf("Current Proc: PID-%d, Level-%d, Set-%d\n", myproc()->pid, myproc()->level, myproc()->set);
            //cprintf("Current RP: Level-%d, Set-%d\n", rp->level, rp->set);
            break;}
    }
    --rp->lvlquantum;
    --myproc()->ticks_left;
```

Figure 6: Code Snippet for Process-Local Quantum Consumption

9. Implementation Explanation for Process-Local Quantum Replenishment

Similarly to the process-local quantum consumption, the replenishment of the quantum process-wise is replenished every time the process is enqueued into a new level as well as every time a process goes to the expired set. Since there are several conditions for replenishment of the process-local quantum, there are several conditionals which check for all these situations. The code snippet regarding this is shown below.

```
if (rp->lvlquantum == 0 || myproc()->ticks_left == 0){  
    if (rp->lvlquantum == 0 && myproc()->level != RSDL_LEVELS - 1 && myproc()->ticks_left == 0){  
        myproc()->ticks_left = RSDL_PROC_QUANTUM;  
        leveldequeue(myproc()->level);  
        yield();  
    }  
  
    else if (rp->lvlquantum == 0 && myproc()->level != RSDL_LEVELS - 1 && myproc()->ticks_left > 0){  
        leveldequeue(myproc()->level);  
        yield();  
    }  
  
    else if (rp->lvlquantum == 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left > 0){  
        leveldequeue(myproc()->level);  
        yield();  
    }  
  
    else if (rp->lvlquantum != 0 && myproc()->level != RSDL_LEVELS - 1 && myproc()->ticks_left == 0){  
        int result = searchlevel(myproc()->level+1);  
        if(result != -1)  
            myproc()->level = result;  
        else{  
            myproc()->set = 1;  
            myproc()->level = myproc()->starting_level;  
        }  
        myproc()->queue = queuesort(myproc()->set, myproc()->level);  
        myproc()->ticks_left = RSDL_PROC_QUANTUM;  
        yield();  
    }  
  
    else if (rp->lvlquantum != 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left == 0){  
        myproc()->set = 1;  
        myproc()->level = myproc()->starting_level;  
        myproc()->queue = queuesort(myproc()->set, myproc()->level);  
        myproc()->ticks_left = RSDL_PROC_QUANTUM;  
        yield();  
    }  
  
    else if (rp->lvlquantum == 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left == 0){  
        myproc()->ticks_left = RSDL_PROC_QUANTUM;  
        leveldequeue(myproc()->level);  
        yield();  
    }  
}
```

Figure 7: Code Snippet Relevant to Process-Local Quantum Replenishment

The codes we are interested in are the conditionals that have the code **myproc()->ticks_left = RSDL_PROC_QUANTUM** which are highlighted in figure 7. The general conditional checks on whether the quantum of the level is 0 since this causes all the processes to get enqueued into a new level thus replenishing their quantum or whether the quantum of the process is emptied out or reaches 0 which causes it either to get enqueued into a lower priority level or into the expired set. The conditional for the first box checks whether the level reaches 0 quantum, whether it's not the lowest priority level, and if the process quantum reaches 0 as well. This condition causes the process to get enqueued into a lower priority level thus replenishing its quantum by setting its quantum (ticks_left) to the original quantum (RSDL_PROC_QUANTUM). The second box checks for the condition wherein the level quantum still exists but the process quantum reaches 0 meaning it needs to be dequeued and enqueued into a lower priority level. Similarly to the first box, it replenishes its quantum by setting the variable of the ticks_left to the original process quantum. Lastly, the third box checks for the condition for the level quantum as well as the process quantum reaching 0 and the level being the lowest priority level in which case it goes to the expired set and replenishes its quantum through the same process.

Additionally, we have also process replenishment for processes who went down a level because of their level-local quantum being exhausted. This is seen in the leveldequeue function and switch_set function. Note that the process in *qq* is the last active process and should be dequeued last and their quantum to be replenished.

```
for(int k = 0; k < NPROC; k++){
    xp = &table.proc[k];
    if(xp->state == RUNNING && xp->level == lvl && xp->set == 0 && xp->state != UNUSED && xp->state != ZOMBIE){
        qq = xp;
        else if(xp->level == lvl && xp->set == 0 && xp->state != UNUSED && xp->state != ZOMBIE){
            if(result != -1){
                xp->level = result;
            }
            else{
                xp->set = 1;
                xp->level = xp->starting_level;
            }
            xp->ticks_left = RSDL_PROC_QUANTUM;
            xp->queue = queuesort(xp->set, xp->level);
        }
    }
}

if(qq){
    if(result != -1){
        qq->level = result;
    }
    else{
        qq->set = 1;
        qq->level = qq->starting_level;
    }
    qq->ticks_left = RSDL_PROC_QUANTUM;
    qq->queue = queuesort(qq->set, qq->level);
}

for (int i = 0; i <= maxqueue; i++){
    for (int k = 0; k < NPROC; k++) {
        pp = &table.proc[k];
        if (pp->set == 0 && (pp->state != UNUSED && pp->state != ZOMBIE) && pp->queue == 1){
            acquire(&table.lock);
            pp->queue = -1;
            pp->set = 1;
            pp->level = pp->starting_level;
            release(&table.lock);
            pp->queue = queuesort(pp->set, pp->level);
        }
    }
}

acquire(&table.lock);
for (int k = 0; k < NPROC; k++) {
    pp = &table.proc[k];
    if (pp->state != UNUSED && pp->state != ZOMBIE){
        pp->ticks_left = RSDL_PROC_QUANTUM;
        pp->set = 0;
    }
}
```

Figure 8: Code Snippets Relevant to Process-Local Quantum Replenishment

10. Implementation Explanation for Schedlog

Similar to lab exercise 5, we were required to make a schedlog syscall. As it is fairly similar, we made the schedlog syscall in the lab exercise as our base code and made changes from thereon. The following steps are very important so that the system will recognize our implemented syscalls when in both user and kernel interface. To implement the schedlog syscall there were several steps.

1. Add `int schedlog(int)` in `user.h`

```
44  int schedlog(int);
```

2. Add `SYSCALL(schedlog)` is `usys.S`

```
34  SYSCALL(schedlog)
```

3. Add `#define SYS schedlog 24` in `syscall.h`

```
25  #define SYS_schedlog 24
```

4. Add `[SYS schedlog] sys schedlog`, and `extern int sys schedlog(void);` in `syscall.c`

```
108  extern int sys_schedlog(void);
```

```
135  [SYS_schedlog] sys_schedlog,
```

5. Add the code block below in `sysproc.c`. The codeblock takes in an integer input (returns -1 if the argument pass has failed) and calls the `schedlog` function in the `proc.c`, passing the maximum number of ticks that the `schedlog` will print.

```
107  int sys_schedlog(void) {  
108      int n;  
109      if(argint(0, &n) < 0)  
110          return -1;  
111      schedlog(n);  
112      return 0;  
113  }
```

6. Add the code block below which forms the `schedlog` function to be used in the scheduler in `proc.c` before the scheduler function. This is used to flag that the `schedlog` is active and tracks how long (in ticks) will the `schedlog` be active.

```
645  void schedlog(int n) {  
646      schedlog_active = 1;  
647      schedlog_lasttick = ticks + n;  
648  }
```

7. Add `void schedlog(int);` in `defs.h`

```
193  void schedlog(int);
```

Recall that although the codes from the lab exercise can mostly be copied, the format for the `schedlog` in the project changes according to the phase. For phase 5, the format is shown below in figure 9.

- **Phase 4 and 5:** <tick>|<set>|<level>(<quantum left>), [<PID>]<process name>:<state number>(<quantum left>),...

Figure 9: Schedlog Format for Phase 5

Recall that the main function of the schedlog syscall is to print out all the relevant information before each context switch. To generate the format shown in figure 9, we adjusted the code provided in lab exercise 5. Below is the relevant code snippet for this format. We separated the printing into two parts, mainly the general information such as the tick, set, and level and process information where we printed down all the necessary processes in that specific set and level as well as its information such as the process-local quanta. The code for printing the general information is shown in figure 8, the 1st part of the code snippet. Notice the conditionals if, else if, and else which are used for printing out the general information for different conditions of the processes. The first conditional is for checking if the process is in the expired set as well as checking if there is a process in the expired set using the function `expired_check`. The second conditional is for checking for active sets and whether there are processes in the active set. The else conditionals is for situations wherein the level has no processes enqueued into them. It's also worth mentioning that the outer loop `i` is for looping through the two sets (active and expired) while the loop for `j` is for looping through the 0 to N-1 levels.

```
if (schedlog_active) {
    if (ticks > schedlog_lasttick) {
        schedlog_active = 0;
    } else {
        for (int i = 0; i < 2; i++){
            for (int j = 0; j < RSDL_LEVELS; j++){
                //Track the highest queue
                struct proc *pp;
                int highest_idx = -1;

                //DEBUG
                /*
                struct qlevel *rrp;
                for(rrp = plevel.qlevel; rrp < &plevel.qlevel[RSDL_LEVELS*2]; rrp++){
                    printf("Set=%d, Level=%d, Quantum=%d\n", rrp->set, rrp->level, rrp->lvlquantum);
                }
                */

                struct qlevel *rp;
                for(rp = plevel.qlevel; rp < &plevel.qlevel[RSDL_LEVELS*2]; rp++){
                    if(rp->level == j && rp->set == i)
                        break;
                }

                if (i == 1 && expired_check(j)) printf("%d|s|d(%d)", ticks, curr_set[i], j, rp->lvlquantum);
                else if (i == 0 && active_check(j)) printf("%d|s|d(%d)", ticks, curr_set[i], j, rp->lvlquantum);
                else printf("%d|s|d(%d)", ticks, curr_set[i], j, rp->lvlquantum);

                for (int k = 0; k < NPROC; k++) {
                    pp = &ptable.proc[k];
                    if (pp->queue > highest_idx && pp->set == i && pp->level == j) {
                        highest_idx = pp->queue;
                    }
                }
            }
        }
    }
}
```

Figure 10: Part 1 of 2 of the Code Snippet for Schedlog

For the second part of the implementation, we can now iterate through the *ptable* as we know the number of current process in the current level and set. We will use this information to print the queue. Meaning the queue should be equal to *highest_idx*. While iterating through the *ptable*, the process with current queue number, its state being unused and adheres to the current set and level will be printed by this loop. This would print the queue processes in order along with their proper information. Lastly, the last for loop looks for the last process in the queue and avoids printing “,” to comply with the formatting stated.

```
for (int x = 0; x < highest_idx; x++){
    for (int k = 0; k <= NPROC; k++){
        pp = &ptable.proc[k];
        if (x == pp->queue && pp->state != UNUSED && pp->set == i && pp->level == j){
            //cprintf("[%d]s:%d(%d), Q=%d set=%d level=%d ", pp->pid, pp->name, pp->state, pp->ticks_left, pp->queue,

            cprintf("[%d]s:%d(%d)", pp->pid, pp->name, pp->state, pp->ticks_left);
            break;
        }
    }
}
for (int k = 0; k <= NPROC; k++) {
    pp = &ptable.proc[k];
    if (highest_idx == pp->queue && pp->state != UNUSED && pp->set == i && pp->level == j && pp->queue >= 0){
        //cprintf("[%d]s:%d(%d) Q=%d set=%d level=%d ", pp->pid, pp->name, pp->state, pp->ticks_left, pp->queue, pp-

        cprintf("[%d]s:%d(%d)", pp->pid, pp->name, pp->state, pp->ticks_left);
        break;
    }
}
```

Figure 11: Part 2 of 2 of the Code Snippet for Schedlog

11. Code Representation of the Expired Set and its Levels

The code representation for the expired set and its levels is found in the `proc.c` file inside the `levelinit` function. This function contains the code representation of both the active and the expired set. The code for this is shown below inside the box. Similar to the active set, we initialize the expired set using the struct `qllevel` which includes the values for the current level (`level`), its set (active or expired using 0 or 1), as well as level quantum (quantum of the specific level). To initialize the expired set, the group used a for loop to loop through all the $N-1$. So as to not overwrite the values of the active level, the group used its initial level for the expired set (0 the highest priority level) as N and set its maximum to $N * 2$. This allowed us to have both the active set and expired set without overwriting their values. For each iteration, we then initialized the values of the level being $i - \text{RSDL_LEVELS}$ since we still need its levels to be labeled as 0 to $N-1$, the value of the set being 1 which corresponds to it being the expired set as well as the quantum level set to the original value which is the `RSDL_LEVEL_QUANTUM`.

```
void levelinit(void){
    struct qllevel *rp;
    //Init active levels
    for(int i=0;i<RSDL_LEVELS;i++){
        rp = &plevel.qllevel[i];
        rp->set = 0;
        rp->level = i;
        rp->lvlquantum = RSDL_LEVEL_QUANTUM;
    }
    //Init expired levels
    for(int i=RSDL_LEVELS;i<RSDL_LEVELS*2;i++){
        rp = &plevel.qllevel[i];
        rp->set = 1;
        rp->level = i - RSDL_LEVELS;
        rp->lvlquantum = RSDL_LEVEL_QUANTUM;
    }
}
```

Figure 12: Code Representation for the Expired Set and its Levels

12. Implementation Explanation for Transferring a Process from Active to Expired

A process will only go to the expired set when its current level is the lowest priority level (N-1) and either the process quantum or the level quantum reaches 0 during which they will then go to the expired set. The implementation for the transfer then starts from the trap.c file where we first check if the quantum reaches 0. This can be seen in the code snippet shown below. Figure 13 shows the condition wherein the level quantum reaches 0 but the current running process still has quantum. When this happens, all the processes from that level will be dequeued and enqueued into the expired set since the current level is $RSDL_LEVELS - 1$ ($N - 1$). Figure 14 shows two conditionals. The first conditional shows when the process in the lowest priority level runs out of quantum wherein only that process will be enqueued into the expired set on the original priority level by calling `leveldequeue`. We also redo the value of the level to the starting level since we want the process to be enqueued into the original priority level. We also call `queuesort()` to make sure that we remove the process in the active queue using `queueclean` as well as to ensure that the process gets enqueued into the last index since we follow the FIFO standard. We also replenish the process' quantum since we have the rule that the quantum gets replenished when we transfer a process from the active to the expired. The last conditional is for when both the process quantum as well as the level quantum reaches 0. For conditionals 1 and 3 we call on `level dequeue` since we dequeue all processes from that specific level whenever our level quantum reaches 0.

```
else if (rp->lvlquantum == 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left > 0){
    leveldequeue(myproc()->level);
    yield();
}
```

Figure 13: First Conditional for Transferring Processes from Active to Expired

```
else if (rp->lvlquantum != 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left == 0){
    myproc()->set = 1;
    myproc()->level = myproc()->starting_level;
    myproc()->queue = queuesort(myproc()->set, myproc()->level);
    myproc()->ticks_left = RSDL_PROC_QUANTUM;
    yield();
}
else if (rp->lvlquantum == 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left == 0){
    myproc()->ticks_left = RSDL_PROC_QUANTUM;
    leveldequeue(myproc()->level);
    yield();
}
```

Figure 14: Conditionals for Transferring Processes from Active to Expired

We then discuss the relevant code in the `leveldequeue` function for which the processes in the lowest priority level all go to the expired set. The relevant code snippet for this is shown below. With the help of the `searchlevel` function, we can determine if there are no available levels left in the active set. In this case, we transfer the processes in that level by setting their set values to 1 keeping in mind that level becomes the original priority level.

```
void leveldequeue(int lvl){
    // Search for processes in the same level, subtract their queue
    // Assign them to another level, Last active goes to the back of the queue
    // Check flag if the level is the last level in the set
    struct proc *xp;
    struct proc *qq;

    int result = searchlevel(lvl);

    for(int k = 0; k < NPROC; k++){
        xp = &ptable.proc[k];
        if(xp->state == RUNNING && xp->level == lvl && xp->set == 0 && xp->state != UNUSED && xp->state != ZOMBIE)
            qq = xp;
        else if(xp->level == lvl && xp->set == 0 && xp->state != UNUSED && xp->state != ZOMBIE){
            if(result != -1)
                xp->level = result;
            else{
                xp->set = 1;
                xp->level = xp->starting_level;
            }
            xp->ticks_left = RSDL_PROC_QUANTUM;
            xp->queue = queuesort(xp->set, xp->level);
        }
    }

    if(qq){
        if(result != -1){
            qq->level = result;
        }
        else{
            qq->set = 1;
            qq->level = qq->starting_level;
        }
        qq->ticks_left = RSDL_PROC_QUANTUM;
        qq->queue = queuesort(qq->set, qq->level);
    }
    return;
}
```

Figure 15: Relevant Code Snippet in leveldequeue Function for Transferring Processes from Active to Expired

13. Implementation Explanation for Swapping Sets

For the implementation of swapping sets, most of the added/modified code is inside the `proc.c` file. First, we need to know when we should swap the sets or to be more specific, we need to know when there are no more processes that are ready to execute in the active set. This happens inside the scheduler function in `proc.c`. There is first a nested for-loop that goes through the queues of each level in the active set looking for the process ready to execute with the lowest queue index. This process is stored to the variable `choice`, however, if no such process is found, then the value of `choice` stays at -1. After the nested for-loop ends, we now enter an if-statement, it checks whether the value of `choice` is either -1 or not. If its value is -1, then there are no more processes that are ready to execute in the active set, which means we shall check if the expired set has a process and if it has, then we shall swap the sets, hence we call the `switch_set` function (more on this later). This happens in a for-loop that calls the `expired_check` function (discussed at the auxiliary functions portion of the document) for each level in the expired set. Of course, once we have confirmed that there is a process in the expired set, we shall break out of the for-loop. Below is a code snippet that contains what has been mentioned so far:

```
667     struct proc *xp;
668     int temp = 2147483647;
669     int choice = -1;
670     for(int j = 0; j < RSDL_LEVELS; j++){
671         for(int k = 0; k < NPROC; k++){
672             xp = &ptable.proc[k];
673             //Find the lowest RUNNABLE process in Queue starting from first level and should not be in the expired set
674             if(xp->state == RUNNABLE && xp->queue < temp && xp->set != 1 && xp->level == j){
675                 temp = xp->queue;
676                 choice = k;
677             }
678         }
679         if(choice != -1) break;
680     }
681
682     //If no RUNNABLE process are available
683     if (choice == -1){
684         // Check if expired set has processes and switch if a process exists.
685         for(int j = 0; j < RSDL_LEVELS; j++){
686             int check = expired_check(j);
687             if (check) {
688                 release(&ptable.lock);
689                 switch_set();
690                 acquire(&ptable.lock);
691                 break;
692             }
693         }
694         goto done;
695     }
696 }
```

Figure 16. Code Snippet for Determining When to Swap Sets

Now, to explain the swapping of sets, we examine the code snippet for the `switch_set` function:

```
537 void switch_set(void){
538     //Keep the queue from the old active set
539     //Set the queue of all proc in the old active to -1
540     //Call queue sort for the new sort in the expired set
541     //Set all the processes set to be active
542     struct proc *pp;
543     struct qllevel *rp;
544     int maxqueue = 0;
545
546     queueclean();
547
548     for (int k = 0; k <= NPROC; k++) {
549         pp = &ptable.proc[k];
550         if((pp->state == SLEEPING || pp->state == RUNNING || pp->state == RUNNABLE) && pp->set == 0){
551             if(pp->queue > maxqueue) maxqueue = pp->queue;
552         }
553     }
554
555     for (int i = 0; i <= maxqueue; i++){
556         for (int k = 0; k < NPROC; k++) {
557             pp = &ptable.proc[k];
558             if (pp->set == 0 && (pp->state != UNUSED && pp->state != ZOMBIE) && pp->queue == i){
559                 acquire(&ptable.lock);
560                 pp->queue = -1;
561                 pp->set = 1;
562                 pp->level = pp->starting_level;
563                 release(&ptable.lock);
564                 pp->queue = queuesort(pp->set, pp->level);
565             }
566         }
567     }
568
569     acquire(&ptable.lock);
570     for (int k = 0; k < NPROC; k++) {
571         pp = &ptable.proc[k];
572         if (pp->state != UNUSED && pp->state != ZOMBIE){
573             pp->ticks_left = RSDL_PROC_QUANTUM;
574             pp->set = 0;
575         }
576     }
577     release(&ptable.lock);
578
579     //Refill the quantum in plevel
580     for(int i=0;i<RSDL_LEVELS*2;i++){
581         rp = &plevel.qllevel[i];
582         rp->lvlquantum = RSDL_LEVEL_QUANTUM;
583     }
```

Figure 17. Code Snippet of `switch_set` Function

We first call the `queueclean` function (explained earlier in the document) to make sure that the queues in each level are in the most recent correct order. The for-loop that starts in line 548 essentially gets the process with the highest index present in any queue

in the old active set. This process's index is stored in the maxqueue variable. Next, for the nested for-loop that starts at line 555, the outer loop iterates from zero to the value in maxqueue to make sure we check all processes in the old active set. Next, we loop through all the processes and check if they are in the old active set and are not an unused or a zombie process. If they satisfy these conditions, then we move them to the expired set and set their level to the starting level. Next, we call the queuesort function for each of these processes so that we are sure that the queue in the new expired set has the correct order.

Now, we are done putting all the processes from the old active set to the old expired set. We now need to put all these processes to the new active set. This happens in the for-loop that starts in line 570, wherein it loops through all the processes and checks if they are not unused and not zombies, and if they are indeed not, we move them to the active set and replenish all their quantum. Note that the processes here are already in a queue so there is no need to check for proper ordering this time.

Lastly, swapping of sets shall replenish the level-local quantum of all the levels. Thus, starting from line 580, we enter a for-loop that iterates through all the levels, both active and expired, and replenishes their quantum.

14. Implementation Explanation for Downgrading of Process Levels

Processes will only go down a level every time its own quantum the process-local quantum is depleted as well as the level-local quantum it's on. The two conditions however are different because the process-local quantum being depleted would mean that only that process will go down a level while depletion of the level-local quantum will make all the process on that level go down a level. The relevant codes for these are found in `trap.c` where we have the conditional on finding out when we need to downgrade a level. The downgrade per level is found however in `proc.c` inside the `leveldequeue` function.

We already discussed the implementation for the transfer of the processes from the active to the expired processes so we will no longer discuss that in this section. We will focus on the downgrading of process levels in the active set. There are 3 conditionals shown in the snippet below that we are interested in. These codes are highlighted by a box and will be discussed one-by-one. The code in the first box is for the condition when both the process-local and the level-local quantum is depleted in which all of the processes in that level will be downgraded using the `leveldequeue` function to be discussed below as well as replenish the quantum of the current process since its quantum got depleted.. The second box takes into account the condition when only the level-local quantum is depleted during which the process that made it get depleted will not replenish its quantum and all of the processes on that level will be downgraded by a level using the `leveldequeue` function. The third condition factors out the scenario wherein only the process-local quantum gets depleted in which only that process will go down a level. During this time, we also replenish its quantum by setting the value of the `ticks_left` variable by the original quantum set in the `rsdl.h`. We downgrade it by a level adding 1 to its current level since 0 is always the highest priority level.

The code we are interested in the `leveldequeue` function is shown in figure 19. How this works is first we loop through all the processes and check if they are in the active set as well as if their process states are viable (i.e. not zombie and unused). We also check if the level they are on is the last level and if they are we don't do anything. The second conditional does the same but instead checks if the process isn't in the last level during which we then downgrade this process to the next level by searching for the appropriate value of level as well as sort out the queue since we will be enqueueing and dequeuing. Note that process who went down a level has their quantum replenish.

```
if (rp->lvlquantum == 0 || myproc()->ticks_left == 0){
    if (rp->lvlquantum == 0 && myproc()->level != RSDL_LEVELS - 1 && myproc()->ticks_left == 0){
        myproc()->ticks_left = RSDL_PROC_QUANTUM;
        leveldequeue(myproc()->level);
        yield();
    }

    else if (rp->lvlquantum == 0 && myproc()->level != RSDL_LEVELS - 1 && myproc()->ticks_left > 0){
        leveldequeue(myproc()->level);
        yield();
    }

    else if (rp->lvlquantum == 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left > 0){
        leveldequeue(myproc()->level);
        yield();
    }

    else if (rp->lvlquantum != 0 && myproc()->level != RSDL_LEVELS - 1 && myproc()->ticks_left == 0){
        int result = searchlevel(myproc()->level+1);
        if(result != -1)
            myproc()->level = result;
        else{
            myproc()->set = 1;
            myproc()->level = myproc()->starting_level;
        }
        myproc()->queue = queuesort(myproc()->set, myproc()->level);
        myproc()->ticks_left = RSDL_PROC_QUANTUM;
        yield();
    }

    else if (rp->lvlquantum != 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left == 0){
        myproc()->set = 1;
        myproc()->level = myproc()->starting_level;
        myproc()->queue = queuesort(myproc()->set, myproc()->level);
        myproc()->ticks_left = RSDL_PROC_QUANTUM;
        yield();
    }

    else if (rp->lvlquantum == 0 && myproc()->level == RSDL_LEVELS - 1 && myproc()->ticks_left == 0){
        myproc()->ticks_left = RSDL_PROC_QUANTUM;
        leveldequeue(myproc()->level);
        yield();
    }
}
```

Figure 18: Code Snippet for Downgrading Process Levels in trap.c

```
int result = searchlevel(lvl);

for(int k = 0; k < NPROC; k++){
    xp = &ptable.proc[k];
    if(xp->state == RUNNING && xp->level == lvl && xp->set == 0 && xp->state != UNUSED && xp->state != ZOMBIE)
        qq = xp;
    else if(xp->level == lvl && xp->set == 0 && xp->state != UNUSED && xp->state != ZOMBIE){
        if(result != -1)
            xp->level = result;
        else{
            xp->set = 1;
            xp->level = xp->starting_level;
        }
        xp->ticks_left = RSDL_PROC_QUANTUM;
        xp->queue = queuesort(xp->set, xp->level);
    }
}

if(qq){
    if(result != -1){
        qq->level = result;
    }
    else{
        qq->set = 1;
        qq->level = qq->starting_level;
    }
    qq->ticks_left = RSDL_PROC_QUANTUM;
    qq->queue = queuesort(qq->set, qq->level);
}
```

Figure 19: Code Snippet for Downgrading Process Levels in proc.c

15. Implementation Explanation for Level-local Quantum Consumption

The implementation for local-level quantum consumption is primarily located in the `trap.c` file. The relevant code for this is shown in figure 20 below. As previously mentioned in the process-local quantum consumption, the program first checks whether there is a process and if its state is set to `RUNNING`. It also checks if the `trapframe` trap number is consistent with the previous user space registers before its context switches. If all of these conditions are satisfied then we know that there is a process currently running. This means that as the process runs, the quantum of the level it is currently at should be decremented / consumed. But first, we need to know which level should have its quantum consumed. Since the local-level quantum is saved in the `qllevel` struct, we need to make sure that the `rp->level` and `rp->set` is consistent with the current process' level and set (hence the for-loop at lines 112-117). Once we have determined the correct level, we can now decrement its level-local quantum every tick by 1. This is done via `--rp->lvlquantum` as seen in line 118.

```
111     if(myproc() && myproc()->state == RUNNING && tf->trapno == T_IRQ0+IRQ_TIMER){
112         for(rp = plevel.qllevel; rp < &plevel.qllevel[RSDL_LEVELS*2]; rp++){
113             if(rp->level == myproc()->level && rp->set == myproc()->set){
114                 //cprintf("Current Proc: PID-%d, Level-%d, Set-%d\n", myproc()->pid, myproc()->level, myproc()->set);
115                 //cprintf("Current RP: Level-%d, Set-%d\n", rp->level, rp->set);
116                 break;}
117         }
118         --rp->lvlquantum;
119         --myproc()->ticks_left;
```

Figure 20: Code Snippet for Level-Local Quantum Consumption

16. Implementation Explanation for Prio-Fork Syscall

The implementation for the Prio-Fork Syscall involves several steps. Firstly, we have to implement the correct system call as stated in the document. This is important so that both user and kernel interfaces will recognize our system call. This is done similarly to the implementation of the *schedlog* syscall. The following are the steps.

Similarly to lab exercise 5, we were required to make a schedlog syscall. As it is fairly similar, we made the schedlog syscall in the lab exercise as our base code and made changes from thereon. The following steps are very important so that the system will recognize our implemented syscalls when in both user and kernel interface. To implement the schedlog syscall there were several steps.

1. Add `int priofork(int)` in `user.h`

```
45  int priofork(int);
```

2. Add `SYSCALL(priofork)` in `usys.S`

```
35  SYSCALL(priofork)
```

3. Add `#define SYS_priofork 25` in `syscall.h`

```
26  #define SYS_priofork 25
```

4. Add `[SYS_priofork] sys_priofork`, and `extern int sys_priofork(void)` in `syscall.c`

```
109  extern int sys_priofork(void);
```

```
136  [SYS_priofork] sys_priofork,
```

5. Add the code block below in `sysproc.c`. The codeblock takes in an integer input (returns -1 if the argument pass has failed) and return the call for priofork function in the `proc.c`, passing the priority level of the forked process.

```
115  int sys_priofork(void){
116      int n;
117      if(argint(0, &n) < 0 || n >= RSDL_LEVELS)
118          return -1;
119      return priofork(n);
120      //return 0;
121  }
```

6. Add `void priofork(int)` in `defs.h`

```
201  int priofork(int);
```

7. Add the code block below in `proc.c`. The codeblock is the implementation of the priofork. The code has copied lines of code from the fork syscall with just a little

modification for both the `starting_level` and `level`. If the `allocproc` succeeds, line 256-257 replaces immediately replaces the process' `level` and `starting_level` to the passed argument respectively. This would effectively set the priority level of the process to whatever the `int` argument is and remembers its `starting_level` whenever we have a set swap. Note that we made changes to the variables compared to the original `fork` function.

```
245 int priofork(int n){
246     int j, pid;
247     struct proc *pfp;
248     struct proc *cur = myproc();
249
250     // Allocate process.
251     if((pfp = allocproc()) == 0){
252         return -1;
253     }
254
255     // Allocate starting level
256     pfp->level = n;
257     pfp->starting_level = n;
258
259     // Copy process state from proc.
260     if((pfp->pgdir = copyuvm(cur->pgdir, cur->sz)) == 0){
261         kfree(pfp->kstack);
262         pfp->kstack = 0;
263         pfp->state = UNUSED;
264         return -1;
265     }
266     pfp->sz = cur->sz;
267     pfp->parent = cur;
268     *pfp->tf = *cur->tf;
269
270     // Clear %eax so that fork returns 0 in the child.
271     pfp->tf->eax = 0;
272
273     for(j = 0; j < NOFILE; j++)
274         if(cur->ofile[j])
275             pfp->ofile[j] = filedup(cur->ofile[j]);
276     pfp->cwd = idup(cur->cwd);
277
278     safestrcpy(pfp->name, cur->name, sizeof(cur->name));
279
280     pid = pfp->pid;
281
282     acquire(&ptable.lock);
283
284     pfp->state = RUNNABLE;
285
286     release(&ptable.lock);
287
288     return pid;
289 }
```

Figure 21: Code Snippet for Prio-Fork Syscall Implementation

VIDEO DOCUMENTATION LINK:

https://drive.google.com/file/d/1lasl4i7iWbzAsCYAY9C4nNdLQkLyQZWq/view?usp=share_link